

```
# Importing libraries

from __future__ import print_function
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.metrics import classification_report
from sklearn import metrics
from sklearn import tree
import warnings
warnings.filterwarnings('ignore')

df = pd.read_csv('/content/crop_recommendation.csv')

df.head()
```

	N	P	K	temperature	humidity	ph	rainfall	label
0	90	42	43	20.879744	82.002744	6.502985	202.935536	rice
1	85	58	41	21.770462	80.319644	7.038096	226.655537	rice
2	60	55	44	23.004459	82.320763	7.840207	263.964248	rice
3	74	35	40	26.491096	80.158363	6.980401	242.864034	rice
4	78	42	42	20.130175	81.604873	7.628473	262.717340	rice

```
df.tail()
```

	N	P	K	temperature	humidity	ph	rainfall	label
2195	107	34	32	26.774637	66.413269	6.780064	177.774507	coffee
2196	99	15	27	27.417112	56.636362	6.086922	127.924610	coffee
2197	118	33	30	24.131797	67.225123	6.362608	173.322839	coffee
2198	117	32	34	26.272418	52.127394	6.758793	127.175293	coffee
2199	104	18	30	23.603016	60.396475	6.779833	140.937041	coffee

```
df.size
```

```
17600
```

```
df.shape
```

```
(2200, 8)
```

```
df.columns
```

```
Index(['N', 'P', 'K', 'temperature', 'humidity', 'ph', 'rainfall', 'label'],  
      dtype='object')
```

```
df['label'].unique()
```

```
array(['rice', 'maize', 'chickpea', 'kidneybeans', 'pigeonpeas',  
      'mothbeans', 'mungbean', 'blackgram', 'lentil', 'pomegranate',  
      'banana', 'mango', 'grapes', 'watermelon', 'muskmelon', 'apple',  
      'orange', 'papaya', 'coconut', 'cotton', 'jute', 'coffee'],  
      dtype=object)
```

```
df.dtypes
```


```
N                int64  
P                int64  
K                int64  
temperature     float64  
humidity        float64  
ph              float64  
rainfall        float64  
label           object  
dtype: object
```

```
df['label'].value_counts()
```

```
rice           100  
maize          100  
jute           100  
cotton         100  
coconut        100  
papaya         100  
orange         100  
apple          100  
muskmelon      100  
watermelon     100  
grapes         100  
mango          100  
banana         100  
pomegranate    100  
lentil         100  
blackgram      100  
mungbean       100  
mothbeans      100  
pigeonpeas     100  
kidneybeans    100  
chickpea       100  
coffee        100  
Name: label, dtype: int64
```

```
sns.heatmap(df.corr(),annot=True)
```

<matplotlib.axes._subplots.AxesSubplot at 0x7f7cab8b7c50>



▼ Separating Features and target label

```
features = df[['N', 'P','K','temperature', 'humidity', 'ph', 'rainfall']]
target = df['label']
#features = df[['temperature', 'humidity', 'ph', 'rainfall']]
labels = df['label']
```

```
# Initializing empty lists to append all model's name and corresponding name
acc = []
model = []
```

```
# Splitting into train and test data
```

```
from sklearn.model_selection import train_test_split
Xtrain, Xtest, Ytrain, Ytest = train_test_split(features,target,test_size = 0.2,random_state
```

▼ Decision Tree

```
from sklearn.tree import DecisionTreeClassifier
```

```
DecisionTree = DecisionTreeClassifier(criterion="entropy",random_state=2,max_depth=5)
```

```
DecisionTree.fit(Xtrain,Ytrain)
```

```

predicted_values = DecisionTree.predict(Xtest)
x = metrics.accuracy_score(Ytest, predicted_values)
acc.append(x)
model.append('Decision Tree')
print("DecisionTrees's Accuracy is: ", x*100)

print(classification_report(Ytest,predicted_values))

```

```

DecisionTrees's Accuracy is: 90.0
      precision    recall  f1-score   support

 apple          1.00      1.00      1.00         13
  banana          1.00      1.00      1.00         17
 blackgram       0.59      1.00      0.74         16
  chickpea          1.00      1.00      1.00         21
   coconut       0.91      1.00      0.95         21
   coffee          1.00      1.00      1.00         22
   cotton          1.00      1.00      1.00         20
   grapes          1.00      1.00      1.00         18
    jute          0.74      0.93      0.83         28
 kidneybeans      0.00      0.00      0.00         14
   lentil         0.68      1.00      0.81         23
   maize          1.00      1.00      1.00         21
   mango          1.00      1.00      1.00         26
  mothbeans       0.00      0.00      0.00         19
  mungbean         1.00      1.00      1.00         24
 muskmelon        1.00      1.00      1.00         23
   orange          1.00      1.00      1.00         29
  papaya          1.00      0.84      0.91         19
 pigeonpeas       0.62      1.00      0.77         18
 pomegranate       1.00      1.00      1.00         17
    rice          1.00      0.62      0.77         16
 watermelon        1.00      1.00      1.00         15

 accuracy                   0.90         440
  macro avg          0.84      0.88      0.85         440
 weighted avg         0.86      0.90      0.87         440

```

```
from sklearn.model_selection import cross_val_score
```

```

# Cross validation score (Decision Tree)
score = cross_val_score(DecisionTree, features, target,cv=5)

```

```
score
```

```
array([0.93636364, 0.90909091, 0.91818182, 0.87045455, 0.93636364])
```

▼ Saving trained Decision Tree Model

```
import pickle
# Dump the trained Naive Bayes classifier with Pickle
DT_pkl_filename = 'DecisionTree.pkl'
# Open the file to save as pkl file
DT_Model_pkl = open(DT_pkl_filename, 'wb')
pickle.dump(DecisionTree, DT_Model_pkl)
# Close the pickle instances
DT_Model_pkl.close()
```

▼ Gaussian Naive Bayes

```
from sklearn.naive_bayes import GaussianNB

NaiveBayes = GaussianNB()

NaiveBayes.fit(Xtrain,Ytrain)

predicted_values = NaiveBayes.predict(Xtest)
x = metrics.accuracy_score(Ytest, predicted_values)
acc.append(x)
model.append('Naive Bayes')
print("Naive Bayes's Accuracy is: ", x)

print(classification_report(Ytest,predicted_values))
```

```
Naive Bayes's Accuracy is: 0.990909090909091
      precision    recall  f1-score   support

   apple         1.00      1.00      1.00        13
  banana         1.00      1.00      1.00        17
 blackgram         1.00      1.00      1.00        16
  chickpea         1.00      1.00      1.00        21
   coconut         1.00      1.00      1.00        21
    coffee         1.00      1.00      1.00        22
    cotton         1.00      1.00      1.00        20
   grapes         1.00      1.00      1.00        18
     jute         0.88      1.00      0.93        28
 kidneybeans       1.00      1.00      1.00        14
    lentil         1.00      1.00      1.00        23
     maize         1.00      1.00      1.00        21
    mango         1.00      1.00      1.00        26
 mothbeans         1.00      1.00      1.00        19
  mungbean         1.00      1.00      1.00        24
 muskmelon         1.00      1.00      1.00        23
   orange         1.00      1.00      1.00        29
  papaya         1.00      1.00      1.00        19
 pigeonpeas       1.00      1.00      1.00        18
 pomegranate       1.00      1.00      1.00        17
     rice         1.00      0.75      0.86        16
 watermelon       1.00      1.00      1.00        15
```

accuracy			0.99	440
macro avg	0.99	0.99	0.99	440
weighted avg	0.99	0.99	0.99	440

```
# Cross validation score (NaiveBayes)
score = cross_val_score(NaiveBayes, features, target, cv=5)
score

array([0.99772727, 0.99545455, 0.99545455, 0.99545455, 0.99090909])
```

▼ Saving Trained Gaussian Bayes Model

```
# Dump the trained Naive Bayes classifier with Pickle
NB_pkl_filename = 'NBClassifier.pkl'
# Open the file to save as pkl file
NB_Model_pkl = open(NB_pkl_filename, 'wb')
pickle.dump(NaiveBayes, NB_Model_pkl)
# Close the pickle instances
NB_Model_pkl.close()
```

▼ Support Vector Machine

```
from sklearn.svm import SVC
# data normalization with sklearn
from sklearn.preprocessing import MinMaxScaler
# fit scaler on training data
norm = MinMaxScaler().fit(Xtrain)
X_train_norm = norm.transform(Xtrain)
# transform testing data
X_test_norm = norm.transform(Xtest)
SVM = SVC(kernel='poly', degree=3, C=1)
SVM.fit(X_train_norm, Ytrain)
predicted_values = SVM.predict(X_test_norm)
x = metrics.accuracy_score(Ytest, predicted_values)
acc.append(x)
model.append('SVM')
print("SVM's Accuracy is: ", x)

print(classification_report(Ytest, predicted_values))

SVM's Accuracy is: 0.9795454545454545
precision    recall  f1-score   support

apple        1.00      1.00      1.00         13
```

banana	1.00	1.00	1.00	17
blackgram	1.00	1.00	1.00	16
chickpea	1.00	1.00	1.00	21
coconut	1.00	1.00	1.00	21
coffee	1.00	0.95	0.98	22
cotton	0.95	1.00	0.98	20
grapes	1.00	1.00	1.00	18
jute	0.83	0.89	0.86	28
kidneybeans	1.00	1.00	1.00	14
lentil	1.00	1.00	1.00	23
maize	1.00	0.95	0.98	21
mango	1.00	1.00	1.00	26
mothbeans	1.00	1.00	1.00	19
mungbean	1.00	1.00	1.00	24
muskmelon	1.00	1.00	1.00	23
orange	1.00	1.00	1.00	29
papaya	1.00	1.00	1.00	19
pigeonpeas	1.00	1.00	1.00	18
pomegranate	1.00	1.00	1.00	17
rice	0.80	0.75	0.77	16
watermelon	1.00	1.00	1.00	15
accuracy			0.98	440
macro avg	0.98	0.98	0.98	440
weighted avg	0.98	0.98	0.98	440

```
# Cross validation score (SVM)
```

```
score = cross_val_score(SVM, features, target, cv=5)
```

```
score
```

```
array([0.97954545, 0.975      , 0.98863636, 0.98863636, 0.98181818])
```

▼ Saving Trained SVM Model

```
# Dump the trained SVM classifier with Pickle
```

```
SVM_pkl_filename = 'SVMClassifier.pkl'
```

```
# Open the file to save as pkl file
```

```
SVM_Model_pkl = open(SVM_pkl_filename, 'wb')
```

```
pickle.dump(SVM, SVM_Model_pkl)
```

```
# Close the pickle instances
```

```
SVM_Model_pkl.close()
```

▼ Logistic Regression

```
from sklearn.linear_model import LogisticRegression
```

```

LogReg = LogisticRegression(random_state=2)

LogReg.fit(Xtrain,Ytrain)

predicted_values = LogReg.predict(Xtest)

x = metrics.accuracy_score(Ytest, predicted_values)
acc.append(x)
model.append('Logistic Regression')
print("Logistic Regression's Accuracy is: ", x)

print(classification_report(Ytest,predicted_values))

```

Logistic Regression's Accuracy is: 0.9522727272727273

	precision	recall	f1-score	support
apple	1.00	1.00	1.00	13
banana	1.00	1.00	1.00	17
blackgram	0.86	0.75	0.80	16
chickpea	1.00	1.00	1.00	21
coconut	1.00	1.00	1.00	21
coffee	1.00	1.00	1.00	22
cotton	0.86	0.90	0.88	20
grapes	1.00	1.00	1.00	18
jute	0.84	0.93	0.88	28
kidneybeans	1.00	1.00	1.00	14
lentil	0.88	1.00	0.94	23
maize	0.90	0.86	0.88	21
mango	0.96	1.00	0.98	26
mothbeans	0.84	0.84	0.84	19
mungbean	1.00	0.96	0.98	24
muskmelon	1.00	1.00	1.00	23
orange	1.00	1.00	1.00	29
papaya	1.00	0.95	0.97	19
pigeonpeas	1.00	1.00	1.00	18
pomegranate	1.00	1.00	1.00	17
rice	0.85	0.69	0.76	16
watermelon	1.00	1.00	1.00	15
accuracy			0.95	440
macro avg	0.95	0.95	0.95	440
weighted avg	0.95	0.95	0.95	440

```

# Cross validation score (Logistic Regression)
score = cross_val_score(LogReg,features,target,cv=5)
score

```

```
array([0.95      , 0.96590909, 0.94772727, 0.96818182, 0.94318182])
```

▼ Saving Trained Logistic Regression Model


```
# Dump the trained Naive Bayes classifier with Pickle
LR_pkl_filename = 'Logistic_Regression.pkl'
# Open the file to save as pkl file
LR_Model_pkl = open(DT_pkl_filename, 'wb')
pickle.dump(LogReg, LR_Model_pkl)
# Close the pickle instances
LR_Model_pkl.close()
```

▼ Random Forest

```
from sklearn.ensemble import RandomForestClassifier

RF = RandomForestClassifier(n_estimators=20, random_state=0)
RF.fit(Xtrain,Ytrain)

predicted_values = RF.predict(Xtest)

x = metrics.accuracy_score(Ytest, predicted_values)
acc.append(x)
model.append('RF')
print("RF's Accuracy is: ", x)

print(classification_report(Ytest,predicted_values))
```

```
RF's Accuracy is: 0.990909090909091
```

	precision	recall	f1-score	support
apple	1.00	1.00	1.00	13
banana	1.00	1.00	1.00	17
blackgram	0.94	1.00	0.97	16
chickpea	1.00	1.00	1.00	21
coconut	1.00	1.00	1.00	21
coffee	1.00	1.00	1.00	22
cotton	1.00	1.00	1.00	20
grapes	1.00	1.00	1.00	18
jute	0.90	1.00	0.95	28
kidneybeans	1.00	1.00	1.00	14
lentil	1.00	1.00	1.00	23
maize	1.00	1.00	1.00	21
mango	1.00	1.00	1.00	26
mothbeans	1.00	0.95	0.97	19
mungbean	1.00	1.00	1.00	24
muskmelon	1.00	1.00	1.00	23
orange	1.00	1.00	1.00	29
papaya	1.00	1.00	1.00	19
pigeonpeas	1.00	1.00	1.00	18
pomegranate	1.00	1.00	1.00	17
rice	1.00	0.81	0.90	16
watermelon	1.00	1.00	1.00	15

accuracy			0.99	440
macro avg	0.99	0.99	0.99	440
weighted avg	0.99	0.99	0.99	440

```
# Cross validation score (Random Forest)
score = cross_val_score(RF, features, target, cv=5)
score

array([0.99772727, 0.99545455, 0.99772727, 0.99318182, 0.98863636])
```

▼ Saving Trained Random Forest Model

```
# Dump the trained Naive Bayes classifier with Pickle
RF_pkl_filename = 'RandomForest.pkl'
# Open the file to save as pkl file
RF_Model_pkl = open(RF_pkl_filename, 'wb')
pickle.dump(RF, RF_Model_pkl)
# Close the pickle instances
RF_Model_pkl.close()
```

▼ XG Boost

```
import xgboost as xgb
XB = xgb.XGBClassifier()
XB.fit(Xtrain, Ytrain)

predicted_values = XB.predict(Xtest)

x = metrics.accuracy_score(Ytest, predicted_values)
acc.append(x)
model.append('XGBoost')
print("XGBoost's Accuracy is: ", x)

print(classification_report(Ytest, predicted_values))
```

```
XGBoost's Accuracy is: 0.9931818181818182
      precision    recall  f1-score   support

   apple         1.00      1.00      1.00         13
  banana         1.00      1.00      1.00         17
 blackgram         1.00      1.00      1.00         16
  chickpea         1.00      1.00      1.00         21
   coconut         1.00      1.00      1.00         21
   coffee         1.00      1.00      1.00         22
```

cotton	1.00	1.00	1.00	20
grapes	1.00	1.00	1.00	18
jute	0.96	0.93	0.95	28
kidneybeans	1.00	1.00	1.00	14
lentil	1.00	1.00	1.00	23
maize	1.00	1.00	1.00	21
mango	1.00	1.00	1.00	26
mothbeans	1.00	1.00	1.00	19
mungbean	1.00	1.00	1.00	24
muskmelon	1.00	1.00	1.00	23
orange	1.00	1.00	1.00	29
papaya	1.00	1.00	1.00	19
pigeonpeas	1.00	1.00	1.00	18
pomegranate	1.00	1.00	1.00	17
rice	0.88	0.94	0.91	16
watermelon	1.00	1.00	1.00	15
accuracy				0.99
macro avg				0.99
weighted avg				0.99

```
# Cross validation score (XGBoost)
score = cross_val_score(XB, features, target, cv=5)
score
```

```
array([0.98636364, 0.99318182, 0.99545455, 0.99090909, 0.98409091])
```


▼ Saving Trained XG Boost Model

```
# Dump the trained Naive Bayes classifier with Pickle
XB_pkl_filename = 'XGBoost.pkl'
# Open the file to save as pkl file
XB_Model_pkl = open(XB_pkl_filename, 'wb')
pickle.dump(XB, XB_Model_pkl)
# Close the pickle instances
XB_Model_pkl.close()
```

▼ Accuracy Comparision

```
plt.figure(figsize=[10,5],dpi = 100)
plt.title('Accuracy Comparison')
plt.xlabel('Accuracy')
plt.ylabel('Algorithm')
sns.barplot(x = acc,y = model,palette='dark')
```

<matplotlib.axes._subplots.AxesSubplot at 0x7f7ca8685550>



```
accuracy_models = dict(zip(model, acc))
for k, v in accuracy_models.items():
    print (k, '-->', v)

Decision Tree --> 0.9
Naive Bayes --> 0.990909090909091
SVM --> 0.9795454545454545
Logistic Regression --> 0.9522727272727273
RF --> 0.990909090909091
XGBoost --> 0.9931818181818182
```

▼ Making a Prediction

```
data = np.array([[104,18, 30, 23.603016, 60.3, 6.7, 140.9]])
prediction = RF.predict(data)
print(prediction)

['coffee']
```

```
data = np.array([[83, 45, 60, 28, 70.3, 7.0, 150.9]])
prediction = RF.predict(data)
print(prediction)
```

```
['jute']
```

[Colab paid products](#) - [Cancel contracts here](#)

